

Sample Questions: MVCC

1 True/False Questions

This document contains sample questions for your quiz on MVCC. On the quiz, each correct answer will result in +1 point, each incorrect answer will result in -1 point, and each blank answer in 0 points.

We will cover the answers to about half the questions in class, and the answers to the other half are from the reading. One of the purposes of the quiz is to test your ability to read/understand documentation, and you will be allowed to use a computer for the quiz to access this documentation.

1. True False An OID is a signed 4-byte integer.
2. True False Every row is assigned an OID.
3. True False The number 35184372088832 is a valid OID.
4. True False The `initdb` utility will initialize a new database cluster.
5. True False The `PGDATA` environment variable stores the path to the base directory.
6. True False It is safe to delete the postgres `pg_xlog` folder in order to free up disk space.
7. True False The file `postgresql.conf` is used to set configuration parameters.
8. True False The folder `pg_stat` contains temporary files for the statistics subsystem.
9. True False Every relation in a postgres database is physically stored on the harddrive as exactly one file.
10. True False With default settings, a table that takes up 160KB on disk has exactly 20 pages.
11. True False TID stands for *Transaction ID*, and every transaction is assigned a unique TID.
12. True False OID stands for *Object ID*, and every table is assigned a unique OID.
13. True False No page can have more than 100 tuples in it.
14. True False Very large tuples can span multiple pages.
15. True False If a tuple exists on disk, then there is guaranteed to be some transaction that can see the tuple.
16. True False In order to remove dead tuples from a table, you must manually run the `VACUUM` command on that table.

- | | | | |
|-----|------|-------|---|
| 17. | True | False | The postgres documentation recommends disabling autovacuum to improve the performance of very large databases. |
| 18. | True | False | The autovacuum will perform a VACUUM FULL operation when needed to reduce the amount of disk space used by a table. |
| 19. | True | False | The <code>pg_freespacemap</code> extension is required to run a VACUUM FULL instead of just a VACUUM. |
| 20. | True | False | Postgres can perform a VACUUM FULL on two different tables at the same time concurrently. |
| 21. | True | False | Postgres can perform a VACUUM on two different tables at the same time concurrently. |
| 22. | True | False | When you use the INSERT command to insert multiple rows into a table at once, these rows are guaranteed to be inserted into the same page. |
| 23. | True | False | When you use the UPDATE command to modify a row, <code>t_xmax</code> field will not be modified. |
| 24. | True | False | Running the DELETE command can modify the <code>t_xmin</code> field of a row. |
| 25. | True | False | A single DELETE command can delete multiple rows from a table. |
| 26. | True | False | A row with <code>t_xmin</code> set to 99 will be visible to a transaction with <code>xid</code> of 90. |
| 27. | True | False | <code>t_cid</code> and <code>t_ctid</code> are synonyms. |
| 28. | True | False | A postgres database cluster can span dozens of computers. |
| 29. | True | False | Postgres is using too much disk space, and you need to free up some space. You identify that there is a large 10TB table that contains about 90% dead tuples. Running the VACUUM command on this table will likely free up several terabytes of disk space. |
| 30. | True | False | Postgres uses a table's FSM to quickly determine which page it can insert a tuple into. |
| 31. | True | False | Postgres uses a table's VM to speed up VACUUMing. |
| 32. | True | False | If you are inserting large text strings into postgres, the postgres documentation recommends compressing them first in order to save space. |
| 33. | True | False | The DELETE command deletes tuples directly from the page files on the harddrive. |
| 34. | True | False | NoSQL databases like MongoDB and CassandraDB use ACID compliant transactions. |
| 35. | True | False | Postgres transactions are ACID compliant. |

- | | | | |
|-----|------|-------|---|
| 36. | True | False | The synchronous_commit system setting can be used to disable postgresql's ACID guarantees. This speeds up transactions, but may result in data loss if the server crashes. |
| 37. | True | False | If the write ahead log (WAL) grows very large, it is safe to delete it in order to free up disk space. |
| 38. | True | False | If the transaction log (clog/xact) grows very large, it is safe to delete it in order to free up disk space. |
| 39. | True | False | If the database cluster is being stored on an SSD, then the random_page_cost system parameter should be reduced from its default value of 4. |
| 40. | True | False | The default fillfactor for tables is 100. |
| 41. | True | False | Increasing a table's fillfactor will generally cause the UPDATE operation to go faster, but the SELECT operation to go slower. |
| 42. | True | False | Tables that have INSERTs but no UPDATEs should use a fillfactor of 100, but for tables with many updates, it is recommended to use a lower fillfactor . |
| 43. | True | False | There is no difference between a line pointer and an item pointer. |
| 44. | True | False | A page always has 24 bytes of header data. |
| 45. | True | False | A page can have 32 bytes of header data if it contains a column with more than 8 non-NULL values. |
| 46. | True | False | For any given page, the pd_lower value can never be greater than the pg_upper value. |
| 47. | True | False | Postgres does not suffer from the txid wraparound problem. |
| 48. | True | False | Phantom reads are possible in Postgres's REPEATABLE READ isolation level. |
| 49. | True | False | The visibility map was introduced in Postgres version 8.4 to reduce the cost of VACUUM processing. |
| 50. | True | False | The visibility map holds information about which pages contain dead tuples. |
| 51. | True | False | The postgres server process is the parent of all other postgres processes. |
| 52. | True | False | The WAL buffer is contained in the shared memory area. |
| 53. | True | False | By default, the maximum number of client connections in postgres is 100. |